

Answers Coding Exercise Module 2

This notebook contains the answers for the coding exercise in Module 2 of the 2025 course “Causal Inference with Linear Regression: A Modern Approach” by CausAI.

Imports

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

import itertools
import random

import statsmodels.api as sm
from scipy import stats

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import KFold
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error

import warnings

warnings.filterwarnings("ignore", category=UserWarning)
```

Question 1

```
data = pd.read_csv("EV_Miles.csv")
data
```

	age	gender	urban_rural	last_year_income
0	62	0	0	68469.13
1	43	1	0	108212.95
2	63	1	0	70245.48
3	48	0	0	72165.09
4	47	1	0	49918.35
...
89858	55	1	0	110483.31

	age	gender	urban_rural	last_year_income
89859	68	1	0	88936.62
89860	43	1	0	98950.97
89861	42	0	0	102328.84
89862	38	0	1	72786.16

	gas_price	owns_ev	miles_driven	fixed_expenses
0	2.97	0	11114	24229.50
1	2.88	1	14432	23737.42
2	3.08	0	14152	10602.99
3	2.98	0	12461	14083.11
4	3.19	0	11246	22264.88
...
89858	3.09	1	14239	22057.39
89859	2.96	0	12715	16675.73
89860	2.81	0	14419	26251.98
89861	3.06	1	11868	27382.49
89862	2.92	0	12471	18386.61

Question 2

```

print("\n--- Shape of Dataset ---")
print(data.shape)

print("\n--- Columns & Data Types ---")
print(data.dtypes)

print("\n--- Unique Values Per Column ---")
for col in data.columns:
    print(f"{col}: {data[col].nunique()} unique values")

data.describe(include="all")

```

--- Shape of Dataset ---
(89863, 8)

--- Columns & Data Types ---
age int64

```

gender          int64
urban_rural    int64
last_year_income float64
gas_price       float64
owns_ev         int64
miles_driven   int64
fixed_expenses float64
dtype: object

--- Unique Values Per Column ---
age: 47 unique values
gender: 2 unique values
urban_rural: 2 unique values
last_year_income: 89282 unique values
gas_price: 95 unique values
owns_ev: 2 unique values
miles_driven: 7045 unique values
fixed_expenses: 87985 unique values

```

	age	gender	urban_rural	last_year_income
count	89863.000000	89863.000000	89863.00000	89863.000000
mean	42.196432	0.496879	0.24899	81093.812888
std	10.143660	0.499993	0.43243	19698.653565
min	24.000000	0.000000	0.00000	31114.510000
25%	34.000000	0.000000	0.00000	66862.575000
50%	42.000000	0.000000	0.00000	78659.890000
75%	49.000000	1.000000	0.00000	92770.785000
max	70.000000	1.000000	1.00000	198228.430000

	gas_price	owns_ev	miles_driven	fixed_expenses
count	89863.000000	89863.000000	89863.000000	89863.000000
mean	3.000364	0.418526	13179.883322	20333.929052
std	0.114016	0.493320	1352.073811	6279.246484
min	2.500000	0.000000	9449.000000	5402.930000
25%	2.920000	0.000000	12238.000000	15778.330000
50%	3.000000	0.000000	13095.000000	19464.570000
75%	3.080000	1.000000	14053.000000	23956.975000
max	3.480000	1.000000	16926.000000	63361.360000

```

for col in data.columns:
    unique_vals = data[col].nunique(dropna=True)
    plt.figure(figsize=(6, 3))

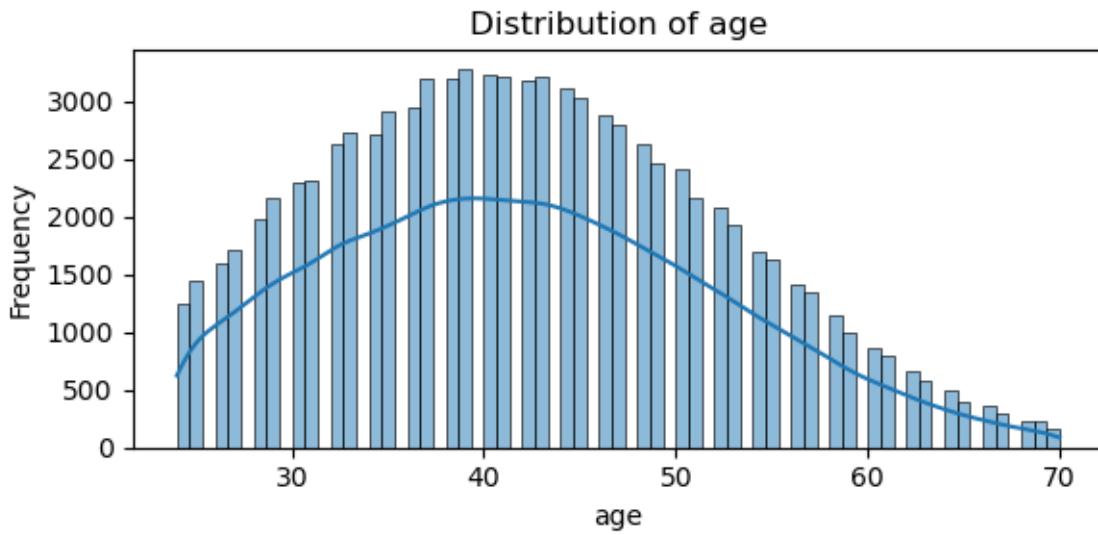
    # Only show KDE if not binary
    sns.histplot(data[col], kde=(unique_vals > 2))

    plt.title(f"Distribution of {col}")
    plt.xlabel(col)
    plt.ylabel("Frequency")
    plt.tight_layout()
    plt.show()

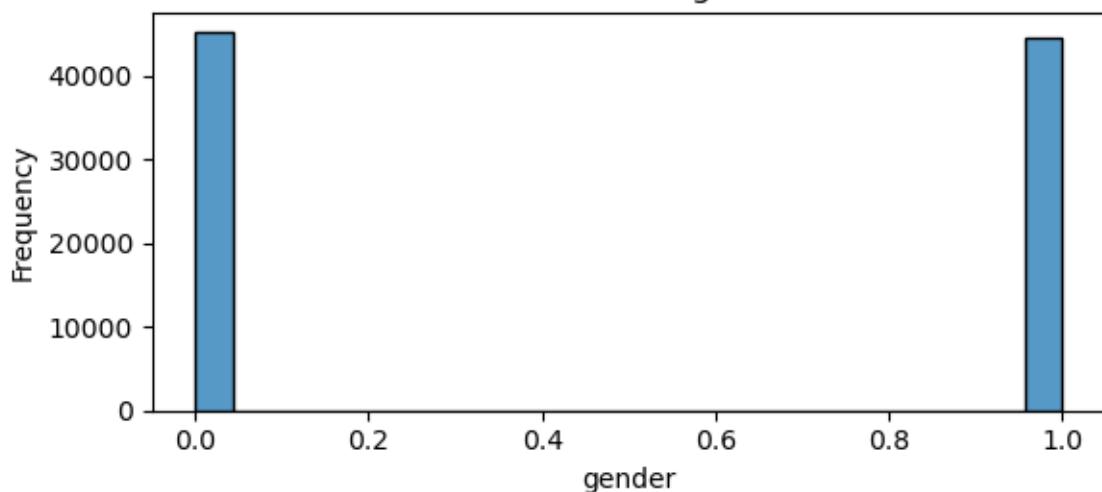
# Box plots to identify outliers (only for non-binary numeric columns)
non_binary_num_cols = [
    col for col in data.columns if data[col].nunique(dropna=True) > 2
]

for col in non_binary_num_cols:
    plt.figure(figsize=(6, 3))
    sns.boxplot(x=data[col])
    plt.title(f"Boxplot of {col}")
    plt.xlabel(col)
    plt.tight_layout()
    plt.show()

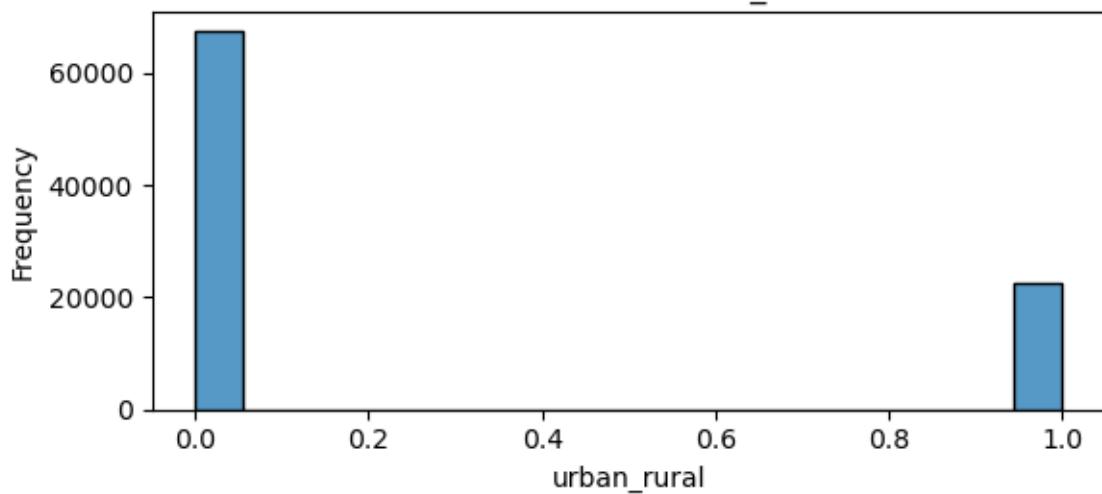
```



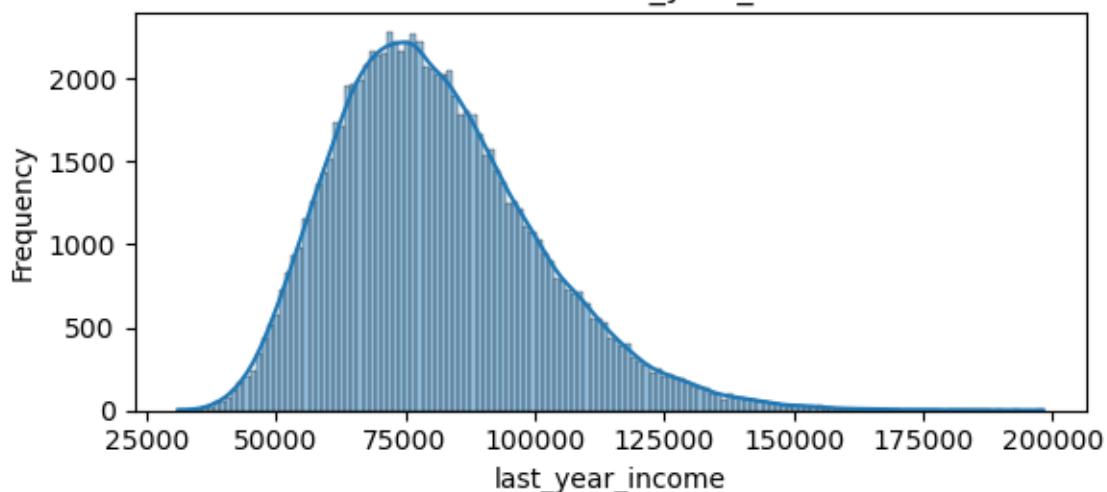
Distribution of gender



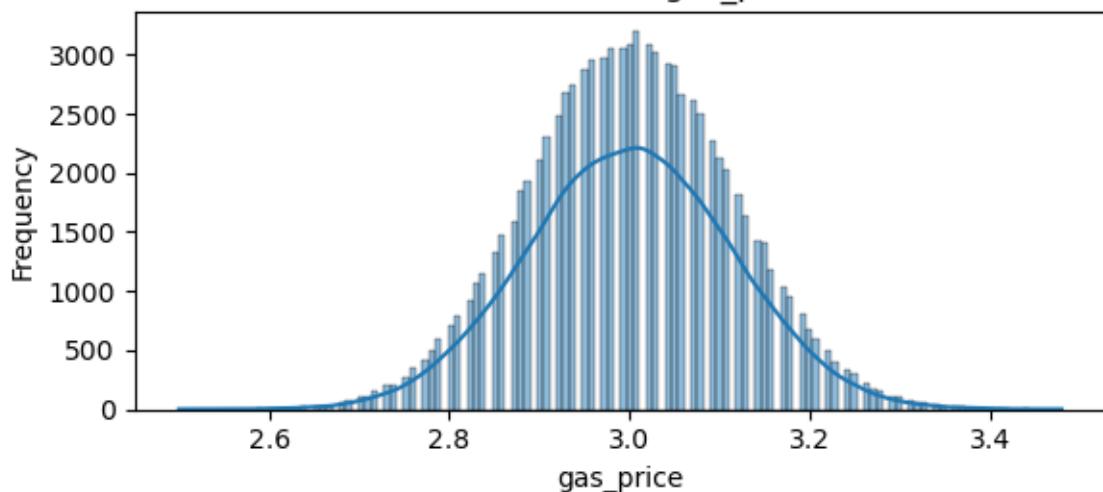
Distribution of urban_rural



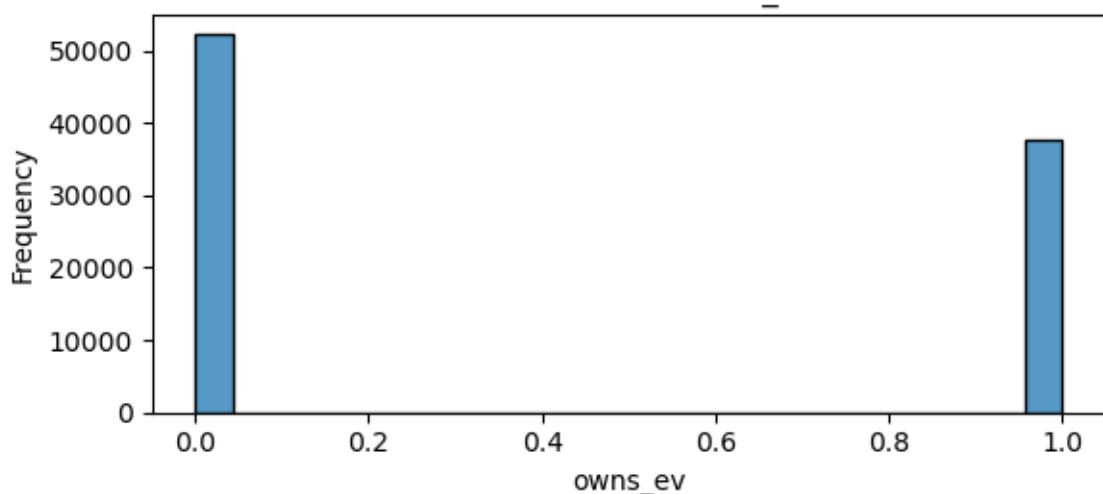
Distribution of last_year_income



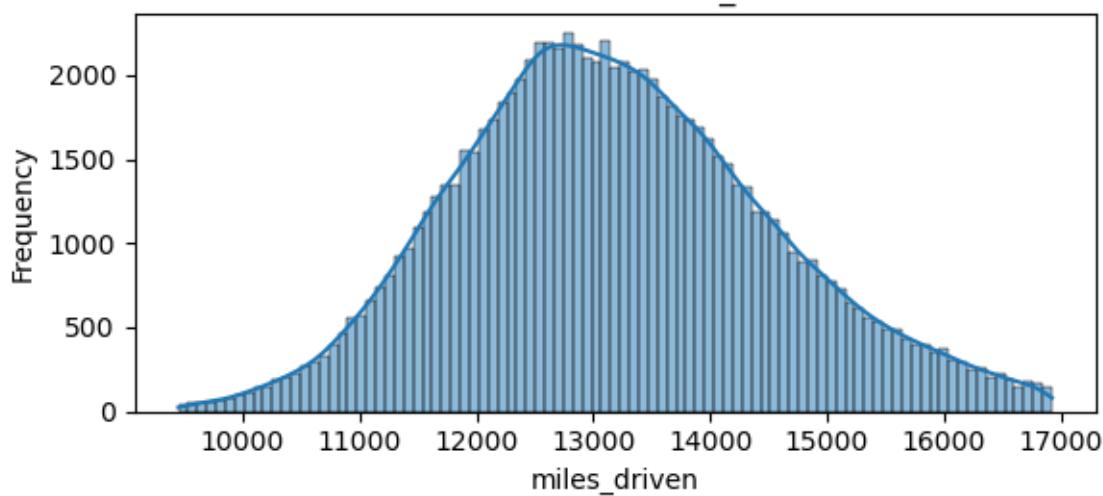
Distribution of gas_price



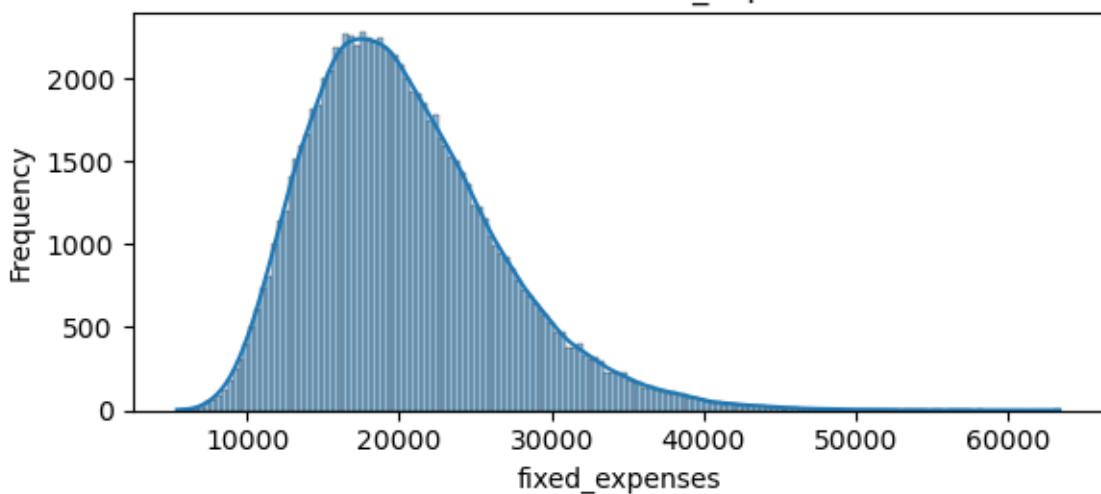
Distribution of `owns_ev`



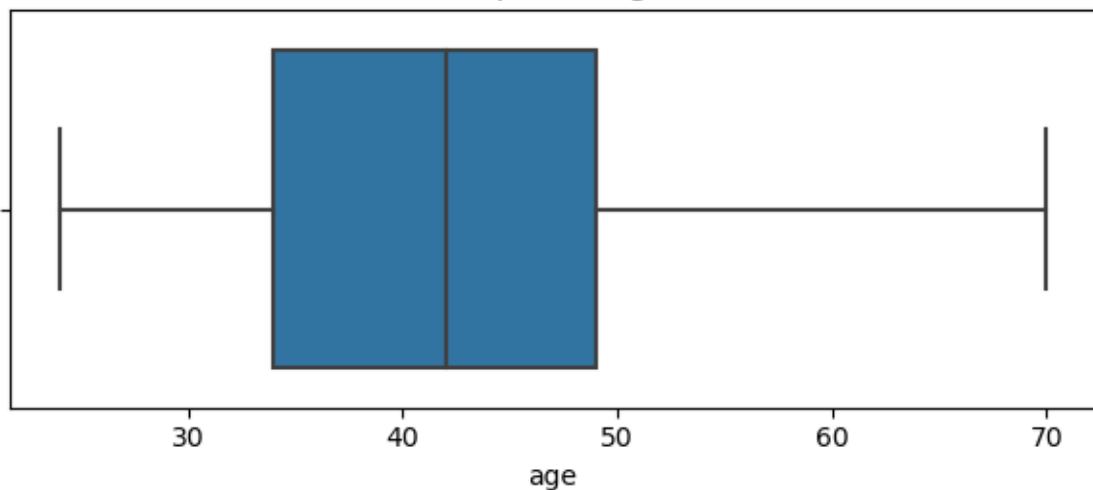
Distribution of `miles_driven`



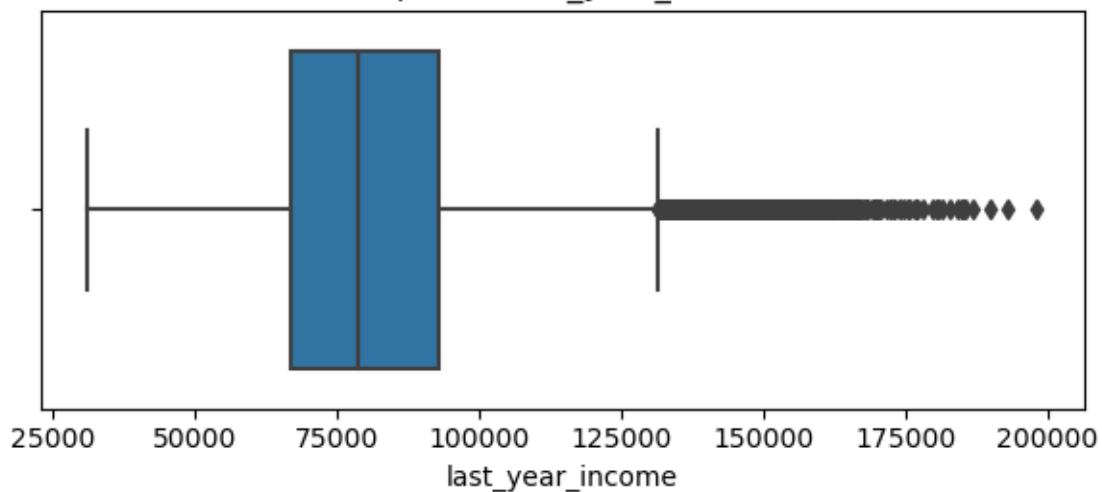
Distribution of fixed_expenses



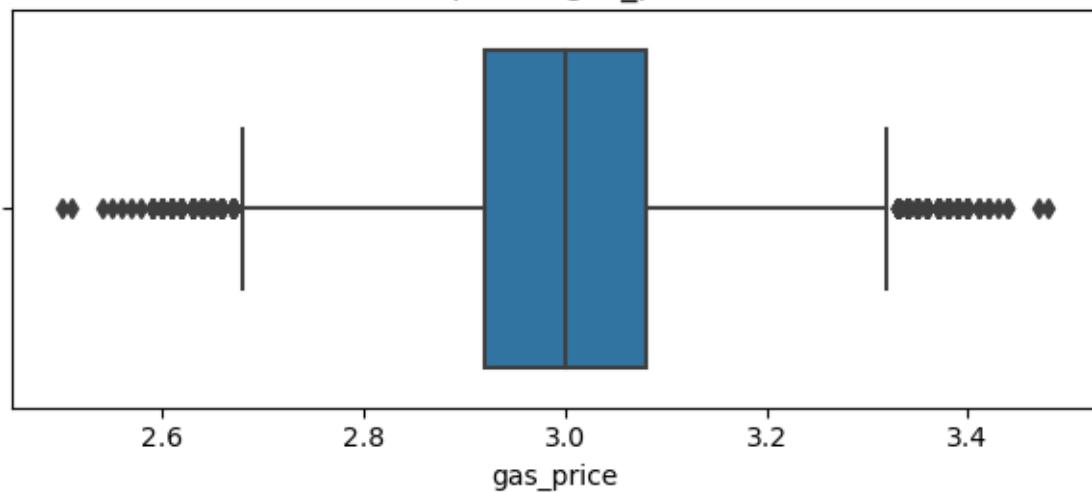
Boxplot of age



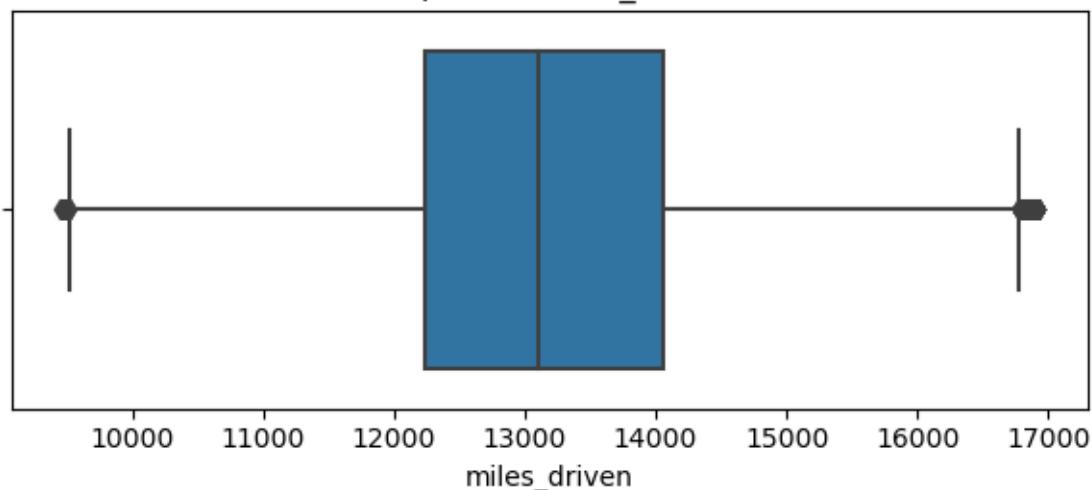
Boxplot of last_year_income



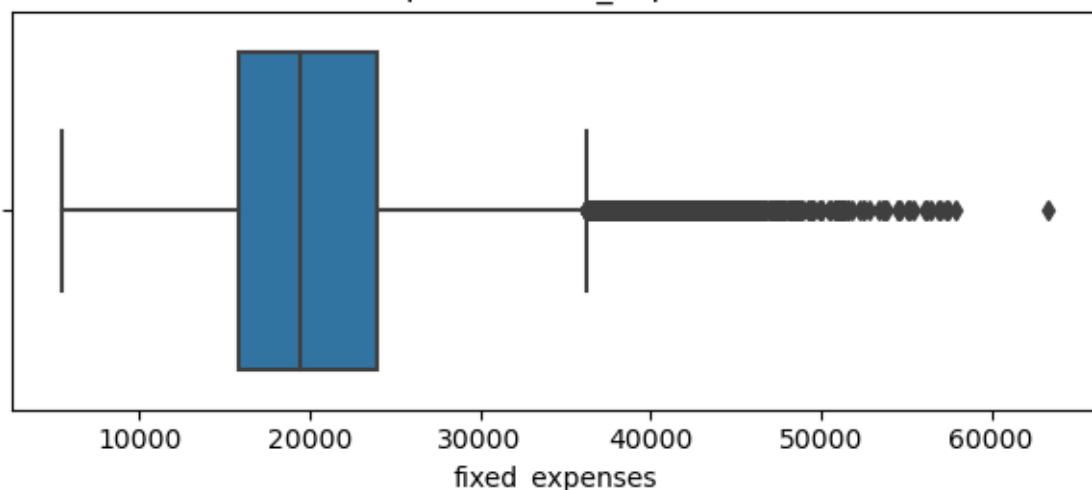
Boxplot of gas_price



Boxplot of miles_driven



Boxplot of fixed_expenses



```
corr_matrix = data.corr()

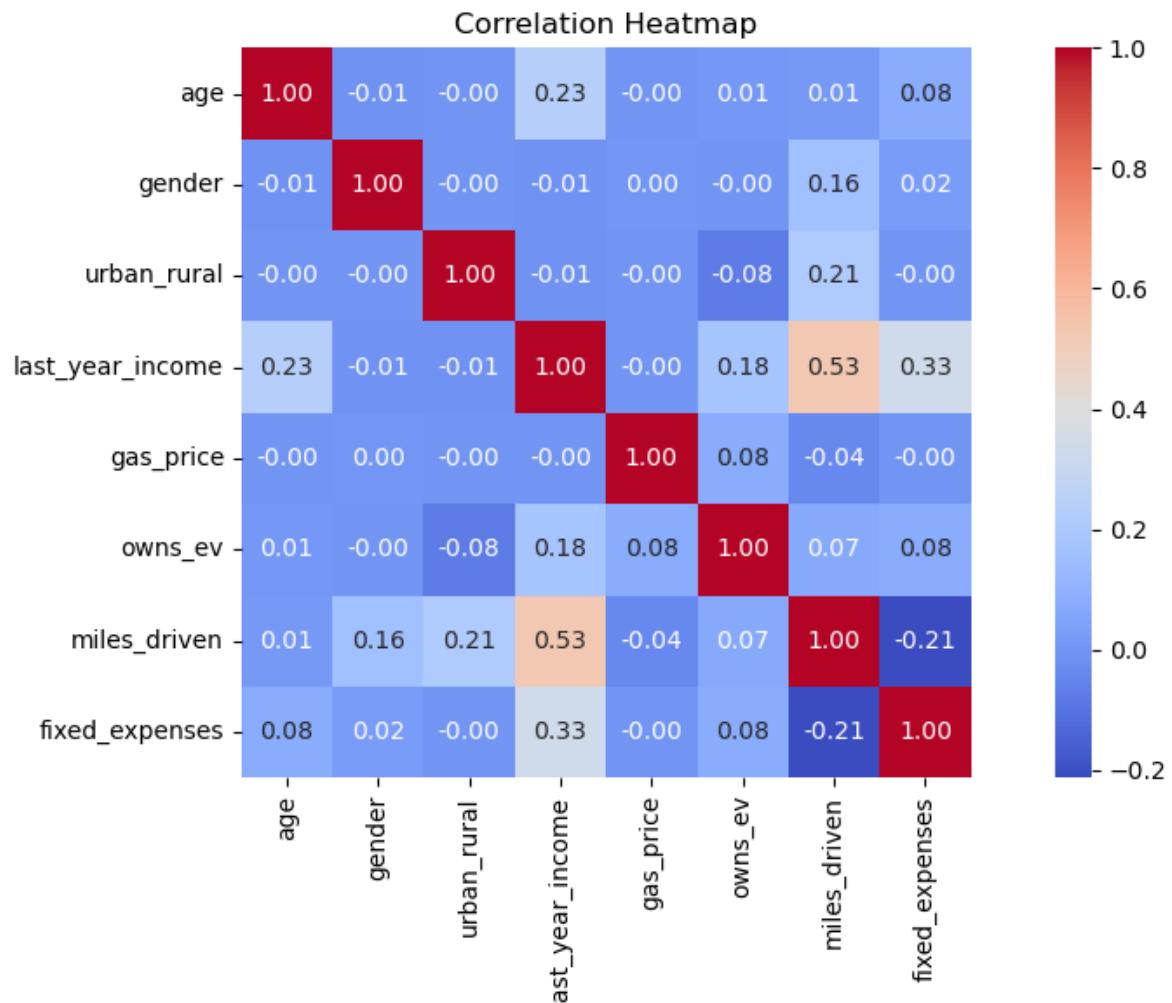
# Heatmap of correlations
plt.figure(figsize=(10, 6))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", fmt=".2f", square=True)
plt.title("Correlation Heatmap")
plt.tight_layout()
plt.show()

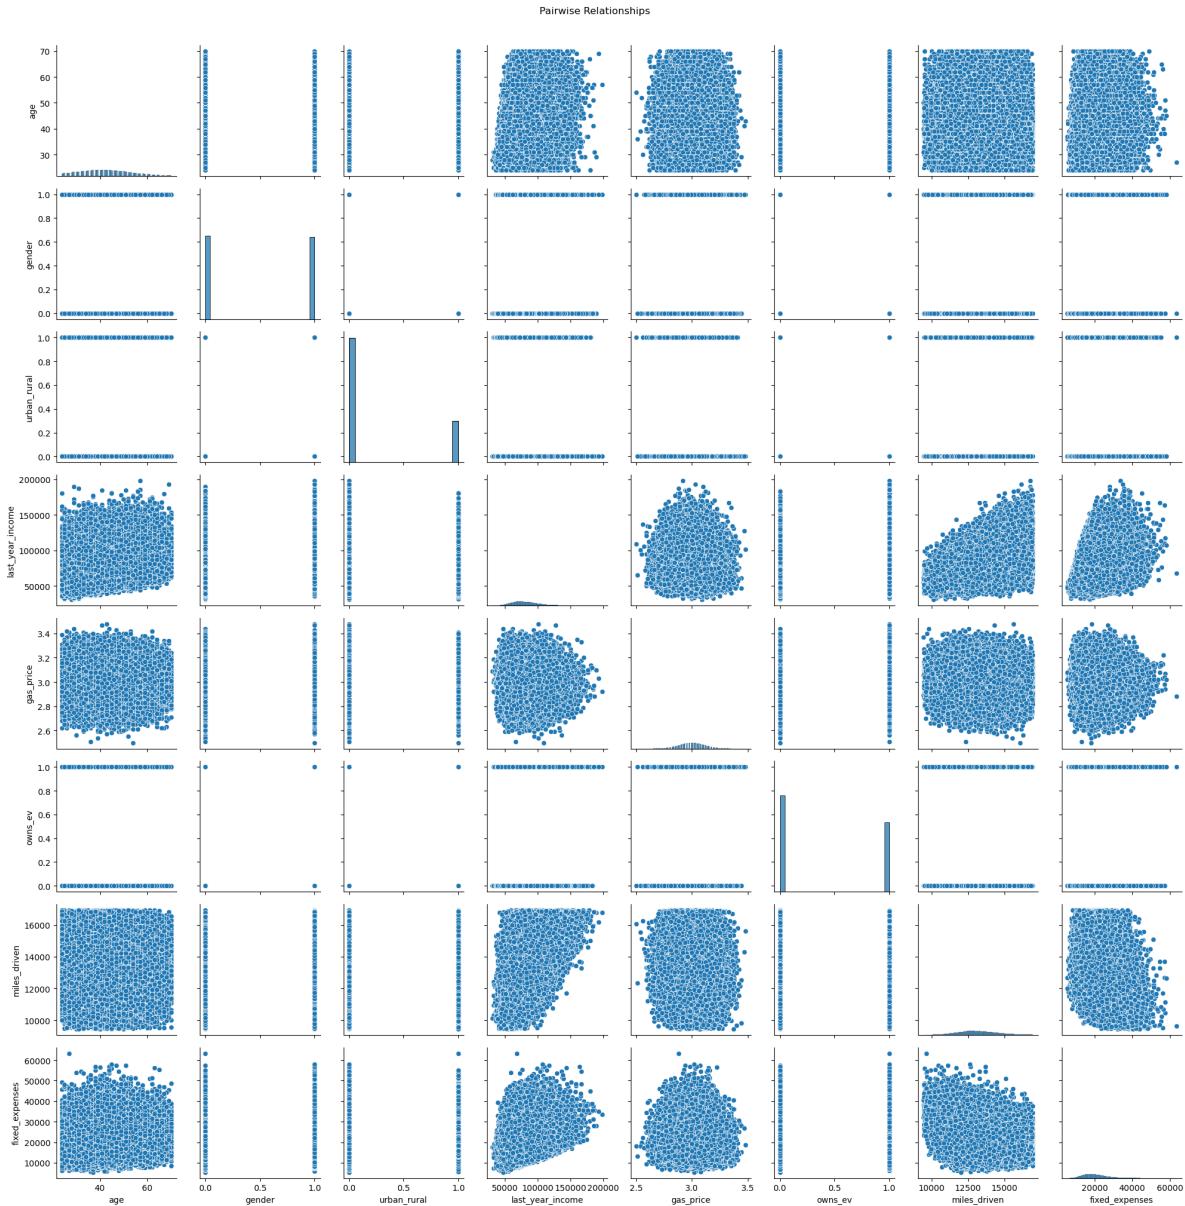
sns.pairplot(data)
```

```

plt.suptitle("Pairwise Relationships", y=1.01)
plt.tight_layout()
plt.show()

```





Question 3

```
mean_difference = np.mean(data["miles_driven"] [data["owns_ev"] == 1]) - np.mean(
    data["miles_driven"] [data["owns_ev"] == 0]
)
mean_difference
```

190.71781349717457

The mean difference is 190.72. This means that on average, owning an EV is associated with 190.72 more annual miles compared to not owning an ev. This difference probably does not reflect the average treatment effect of owning an EV on annual miles driven. There are likely several common causes of owning an EV and miles driven, for example gas prices, income, or age.

Question 4

I picked some arbitrary sets of variables

```
# All variables we can choose to add to the model
all_vars = [
    "age",
    "gender",
    "urban_rural",
    "last_year_income",
    "gas_price",
    "fixed_expenses",
]

# Standardize non-binary variables for improved numerical stability
to_standardize = ["age", "last_year_income", "gas_price", "fixed_expenses"]

data_scaled = data.copy()
scaler = StandardScaler()
data_scaled[to_standardize] = scaler.fit_transform(data[to_standardize])

# Generate random adjustment sets
random.seed(3)

all_combos = []

for r in range(2, len(all_vars) + 1):
    combos = list(itertools.combinations(all_vars, r))
    all_combos.extend(combos)

# Shuffle and take 20 unique random ones
random.shuffle(all_combos)
selected_combos = all_combos[:20]

# Run regressions and print owns_ev coefficient
for i, adjustment_set in enumerate(selected_combos, 1):
    predictors = ["owns_ev"] + list(adjustment_set)
    X = data_scaled[predictors]
```

```

y = data_scaled["miles_driven"]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()

coeff = model.params["owns_ev"]
std_err = model.bse["owns_ev"]
p_val = model.pvalues["owns_ev"]

print(f"\nmiles_driven ~ owns_ev + {' , '.join(adjustment_set)}")
print(
    f"owns_ev coefficient: {coeff:.4f} | Std Err: {std_err:.4f} | p-value: {p_val:.4f}"
)

```

miles_driven ~ owns_ev + urban_rural, gas_price, fixed_expenses
 owns_ev coefficient: 296.0122 | Std Err: 8.7389 | p-value: 0.0000

miles_driven ~ owns_ev + gender, urban_rural, gas_price, fixed_expenses
 owns_ev coefficient: 297.2670 | Std Err: 8.6078 | p-value: 0.0000

miles_driven ~ owns_ev + age, urban_rural, gas_price, fixed_expenses
 owns_ev coefficient: 296.0858 | Std Err: 8.7351 | p-value: 0.0000

miles_driven ~ owns_ev + urban_rural, gas_price
 owns_ev coefficient: 247.6860 | Std Err: 8.9443 | p-value: 0.0000

miles_driven ~ owns_ev + gender, last_year_income
 owns_ev coefficient: -67.0465 | Std Err: 7.7449 | p-value: 0.0000

miles_driven ~ owns_ev + age, gas_price
 owns_ev coefficient: 200.6857 | Std Err: 9.1371 | p-value: 0.0000

miles_driven ~ owns_ev + age, urban_rural, last_year_income
 owns_ev coefficient: -32.1564 | Std Err: 7.5941 | p-value: 0.0000

miles_driven ~ owns_ev + age, urban_rural, gas_price
 owns_ev coefficient: 247.5256 | Std Err: 8.9439 | p-value: 0.0000

miles_driven ~ owns_ev + age, last_year_income, gas_price
 owns_ev coefficient: -70.3397 | Std Err: 7.8418 | p-value: 0.0000

```
miles_driven ~ owns_ev + gender, last_year_income, gas_price, fixed_expenses
owns_ev coefficient: -33.6449 | Std Err: 6.7170 | p-value: 0.0000
```

```
miles_driven ~ owns_ev + gender, urban_rural, gas_price
owns_ev coefficient: 248.1662 | Std Err: 8.8244 | p-value: 0.0000
```

```
miles_driven ~ owns_ev + gender, gas_price
owns_ev coefficient: 201.3084 | Std Err: 9.0208 | p-value: 0.0000
```

```
miles_driven ~ owns_ev + gender, last_year_income, gas_price
owns_ev coefficient: -58.4921 | Std Err: 7.7607 | p-value: 0.0000
```

```
miles_driven ~ owns_ev + age, gender, urban_rural, gas_price
owns_ev coefficient: 247.9900 | Std Err: 8.8238 | p-value: 0.0000
```

```
miles_driven ~ owns_ev + urban_rural, fixed_expenses
owns_ev coefficient: 285.2038 | Std Err: 8.7251 | p-value: 0.0000
```

```
miles_driven ~ owns_ev + age, gender, gas_price
owns_ev coefficient: 201.1345 | Std Err: 9.0204 | p-value: 0.0000
```

```
miles_driven ~ owns_ev + last_year_income, fixed_expenses
owns_ev coefficient: -42.3215 | Std Err: 6.9028 | p-value: 0.0000
```

```
miles_driven ~ owns_ev + urban_rural, last_year_income, gas_price, fixed_expenses
owns_ev coefficient: 13.5017 | Std Err: 6.6399 | p-value: 0.0420
```

```
miles_driven ~ owns_ev + age, gender, urban_rural, last_year_income, gas_price, fixed_expenses
owns_ev coefficient: 0.9733 | Std Err: 6.3469 | p-value: 0.8781
```

```
miles_driven ~ owns_ev + age, gender, last_year_income, gas_price, fixed_expenses
owns_ev coefficient: -45.9296 | Std Err: 6.6338 | p-value: 0.0000
```

In each regression, you can interpret the coefficient for *owns_ev* as “Holding all other variables in the model constant, owning an EV instead of a gasoline car is linearly associated with (coefficient) more annual miles, on average”. You can clearly see that this coefficient varies widely among the different regressions, from negative values like -58 to positive values like 297.

Question 5

```

# Cross-validation setup
kf = KFold(n_splits=5, shuffle=True, random_state=3)
results = []

# Evaluate each 4-variable combination
 combos = list(itertools.combinations(all_vars, 4))

for combo in combos:
    predictors = ["owns_ev"] + list(combo)
    X = data_scaled[predictors].values
    y = data_scaled["miles_driven"].values

    mse_scores = []

    for train_idx, test_idx in kf.split(X):
        X_train, X_test = X[train_idx], X[test_idx]
        y_train, y_test = y[train_idx], y[test_idx]

        model = LinearRegression()
        model.fit(X_train, y_train)
        preds = model.predict(X_test)
        mse = mean_squared_error(y_test, preds)
        mse_scores.append(mse)

    avg_mse = np.mean(mse_scores)
    results.append({"combo": combo, "avg_mse": avg_mse})

# Find the best model
best_model = min(results, key=lambda x: x["avg_mse"])

# Output
print("\nBest 4-variable adjustment set (with owns_ev):\n")
print("Combo:", best_model["combo"])
print("Average CV MSE:", round(best_model["avg_mse"], 2))

```

Best 4-variable adjustment set (with owns_ev):

Combo: ('gender', 'urban_rural', 'last_year_income', 'fixed_expenses')
 Average CV MSE: 868677.37

```

# Use the best combo from previous step
optimal_combo = best_model["combo"]
predictors = ["owns_ev"] + list(optimal_combo)

# Prepare X and y
X = data_scaled[predictors]
y = data_scaled["miles_driven"]

# Add intercept
X = sm.add_constant(X)

# Fit the final model
final_model = sm.OLS(y, X).fit()

# Print summary
print("\nFinal Regression Model Summary (Best variable set in terms of MSE):\n")
print(f"Independent Variables: owns_ev, {''.join(optimal_combo)}\n")
print(final_model.summary())

```

Final Regression Model Summary (Best variable set in terms of MSE):

Independent Variables: owns_ev, gender, urban_rural, last_year_income, fixed_expenses

OLS Regression Results

Dep. Variable:	miles_driven	R-squared:	0.525			
Model:	OLS	Adj. R-squared:	0.525			
Method:	Least Squares	F-statistic:	1.985e+04			
Date:	Thu, 27 Mar 2025	Prob (F-statistic):	0.00			
Time:	14:02:13	Log-Likelihood:	-7.4193e+05			
No. Observations:	89863	AIC:	1.484e+06			
Df Residuals:	89857	BIC:	1.484e+06			
Df Model:	5					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	1.277e+04	5.516	2315.978	0.000	1.28e+04	1.28e+04
owns_ev	4.1864	6.424	0.652	0.515	-8.405	16.778
gender	475.9527	6.221	76.510	0.000	463.760	488.145
urban_rural	671.8651	7.212	93.161	0.000	657.730	686.000

```

last_year_income    916.8499      3.343     274.245      0.000     910.297     923.402
fixed_expenses     -599.3389      3.301    -181.553      0.000    -605.809    -592.869
=====
Omnibus:           17197.669   Durbin-Watson:          1.995
Prob(Omnibus):     0.000     Jarque-Bera (JB):       32678.967
Skew:              1.184     Prob(JB):                  0.00
Kurtosis:          4.767     Cond. No.                 3.38
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We can interpret the coefficient for *owns_ev* in this model as follows: “For individuals with the same gender, living location (urban vs rural), last year’s income and fixed expenses, owning an EV is linearly associated with 4.19 more annual miles compared to not owning an EV, on average”

Question 6

```

# Cross-validation setup
kf = KFold(n_splits=5, shuffle=True, random_state=3)
results = []

# Evaluate each 5-variable combination

 combos = list(itertools.combinations(all_vars, 5))

for combo in combos:
    predictors = ["owns_ev"] + list(combo)
    X = data_scaled[predictors].values
    y = data_scaled["miles_driven"].values

    mse_scores = []

    for train_idx, test_idx in kf.split(X):
        X_train, X_test = X[train_idx], X[test_idx]
        y_train, y_test = y[train_idx], y[test_idx]

        model = LinearRegression()
        model.fit(X_train, y_train)
        preds = model.predict(X_test)
        mse = mean_squared_error(y_test, preds)
        mse_scores.append(mse)
```

```

avg_mse = np.mean(mse_scores)
results.append({"combo": combo, "avg_mse": avg_mse})

# Find the best model
best_model = min(results, key=lambda x: x["avg_mse"])

# Output
print("\nBest 5-variable adjustment set (with owns_ev):\n")
print("Combo:", best_model["combo"])
print("Average CV MSE:", round(best_model["avg_mse"], 2))

```

Best 5-variable adjustment set (with owns_ev):

Combo: ('age', 'gender', 'urban_rural', 'last_year_income', 'fixed_expenses')
 Average CV MSE: 844479.02

```

# Use the best combo from previous step
optimal_combo = best_model["combo"]
predictors = ["owns_ev"] + list(optimal_combo)

# Prepare X and y
X = data_scaled[predictors]
y = data_scaled["miles_driven"]

# Add intercept
X = sm.add_constant(X)

# Fit the final model
final_model = sm.OLS(y, X).fit()

# Print summary
print("\nFinal Regression Model Summary (Best variable set in terms of MSE):\n")
print(f"Independent Variables: owns_ev, {''.join(optimal_combo)}\n")
print(final_model.summary())

```

Final Regression Model Summary (Best variable set in terms of MSE):

Independent Variables: owns_ev, age, gender, urban_rural, last_year_income, fixed_expenses

OLS Regression Results

Dep. Variable:	miles_driven	R-squared:	0.538			
Model:	OLS	Adj. R-squared:	0.538			
Method:	Least Squares	F-statistic:	1.745e+04			
Date:	Thu, 27 Mar 2025	Prob (F-statistic):	0.00			
Time:	14:02:14	Log-Likelihood:	-7.4066e+05			
No. Observations:	89863	AIC:	1.481e+06			
Df Residuals:	89856	BIC:	1.481e+06			
Df Model:	6					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	1.278e+04	5.440	2349.500	0.000	1.28e+04	1.28e+04
owns_ev	-8.0784	6.339	-1.274	0.203	-20.502	4.345
age	-160.2235	3.156	-50.774	0.000	-166.408	-154.038
gender	474.5957	6.134	77.378	0.000	462.574	486.617
urban_rural	669.8451	7.111	94.201	0.000	655.908	683.782
last_year_income	955.5196	3.383	282.440	0.000	948.889	962.150
fixed_expenses	-599.4554	3.255	-184.174	0.000	-605.835	-593.076
Omnibus:	18700.899	Durbin-Watson:	1.994			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	37398.372			
Skew:	1.254	Prob(JB):	0.00			
Kurtosis:	4.924	Cond. No.	3.38			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We can interpret the coefficient for *owns_ev* in this model as follows: “For individuals with the same gender, living location (urban vs rural), last year’s income, fixed expenses and age, owning an EV is linearly associated with 8.08 less annual miles compared to not owning an EV, on average”

Question 7

No, we can’t. Linear regression coefficients are purely measures of association, and we have no idea whether these measured associations are biased or not (they likely are) or to which extent. And although the final models are optimal for predictive accuracy, this doesn’t necessarily mean that they are also optimal for causal inference.

This notebook was converted with convert.ploomber.io